

# prediction-project-codsoft-task-1

October 30, 2023

## Titanic Survival Prediction Project-Giriraju B

Import Data from csv file

```
[223]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, \
    f1_score, \
confusion_matrix, ConfusionMatrixDisplay, RocCurveDisplay
```

```
[224]: df=pd.read_csv('/content/Titanic Data set.csv')
```

```
[225]: df.head()
```

```
[225]: PassengerId  Survived  Pclass  \
0            892         0         3
1            893         1         3
2            894         0         2
3            895         0         3
4            896         1         3
```

```

                                Name    Sex  Age  SibSp  Parch  \
0                      Kelly, Mr. James  male  34.5    0     0
1      Wilkes, Mrs. James (Ellen Needs)  female  47.0    1     0
2                Myles, Mr. Thomas Francis  male  62.0    0     0
3                      Wirz, Mr. Albert  male  27.0    0     0
4  Hirvonen, Mrs. Alexander (Helga E Lindqvist)  female  22.0    1     1
```

```

Ticket    Fare  Cabin  Embarked
0  330911   7.8292   NaN         Q
1  363272   7.0000   NaN         S
2  240276   9.6875   NaN         Q
3  315154   8.6625   NaN         S
4  3101298  12.2875   NaN         S
```

```
[226]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 418 entries, 0 to 417
Data columns (total 12 columns):
 #   Column        Non-Null Count  Dtype
---  -
 0   PassengerId   418 non-null    int64
 1   Survived      418 non-null    int64
 2   Pclass        418 non-null    int64
 3   Name          418 non-null    object
 4   Sex           418 non-null    object
 5   Age           332 non-null    float64
 6   SibSp         418 non-null    int64
 7   Parch         418 non-null    int64
 8   Ticket        418 non-null    object
 9   Fare          417 non-null    float64
10   Cabin         91 non-null     object
11   Embarked      418 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 39.3+ KB
```

```
[227]: df.isnull().sum()
```

```
[227]: PassengerId    0
      Survived      0
      Pclass        0
      Name          0
      Sex           0
      Age           86
      SibSp         0
      Parch         0
      Ticket        0
      Fare          1
      Cabin        327
      Embarked      0
      dtype: int64
```

Data wrangling

```
[228]: df['Age'].fillna(df['Age'].mean(), inplace=True)
```

```
[229]: df['Fare'].fillna(df['Fare'].mean(), inplace=True)
```

```
[230]: df['Embarked'].fillna(df['Embarked'].mode()[0], inplace=True)
```

```
[231]: df.drop(columns=['Cabin'], inplace=True)
```

```
[232]: df.isnull().sum()
```

```
[232]: PassengerId    0
      Survived      0
      Pclass       0
      Name         0
      Sex          0
      Age          0
      SibSp        0
      Parch        0
      Ticket       0
      Fare         0
      Embarked     0
      dtype: int64
```

```
[233]: df.describe()
```

```
[233]:
```

	PassengerId	Survived	Pclass	Age	SibSp \
count	418.000000	418.000000	418.000000	418.000000	418.000000
mean	1100.500000	0.363636	2.265550	30.272590	0.447368
std	120.810458	0.481622	0.841838	12.634534	0.896760
min	892.000000	0.000000	1.000000	0.170000	0.000000
25%	996.250000	0.000000	1.000000	23.000000	0.000000
50%	1100.500000	0.000000	3.000000	30.272590	0.000000
75%	1204.750000	1.000000	3.000000	35.750000	1.000000
max	1309.000000	1.000000	3.000000	76.000000	8.000000

	Parch	Fare
count	418.000000	418.000000
mean	0.392344	35.627188
std	0.981429	55.840500
min	0.000000	0.000000
25%	0.000000	7.895800
50%	0.000000	14.454200
75%	0.000000	31.500000
max	9.000000	512.329200

```
[234]: df['Survived'].value_counts()
```

```
[234]: 0    266
      1    152
      Name: Survived, dtype: int64
```

```
[235]: df['Sex'].value_counts()
```

```
[235]: male    266
      female  152
```

Name: Sex, dtype: int64

So we can understand all the males died on saving females,lol!!


```
[236]: df['Embarked'].value_counts()
```

```
[236]: S    270  
      C    102  
      Q     46  
      Name: Embarked, dtype: int64
```

### Data visualization

```
[237]: sns.set()
```

```
[238]: sns.countplot(x='Sex',data=df)  
plt.title('Gender graph')  
  
ax = sns.countplot(x='Sex', data=df)  
  
total = len(df['Sex'])  
for p in ax.patches:  
    count = p.get_height()  
    percentage = 100 * count / total  
    x = p.get_x() + p.get_width() / 2  
    y = p.get_height() + 0.05  
    ax.annotate(f'{int(count)}({percentage:.2f}%)', (x, y), ha='center',  
↪fontsize=10, color='black')  
  
plt.title('Gender graph')  
  
plt.show()
```

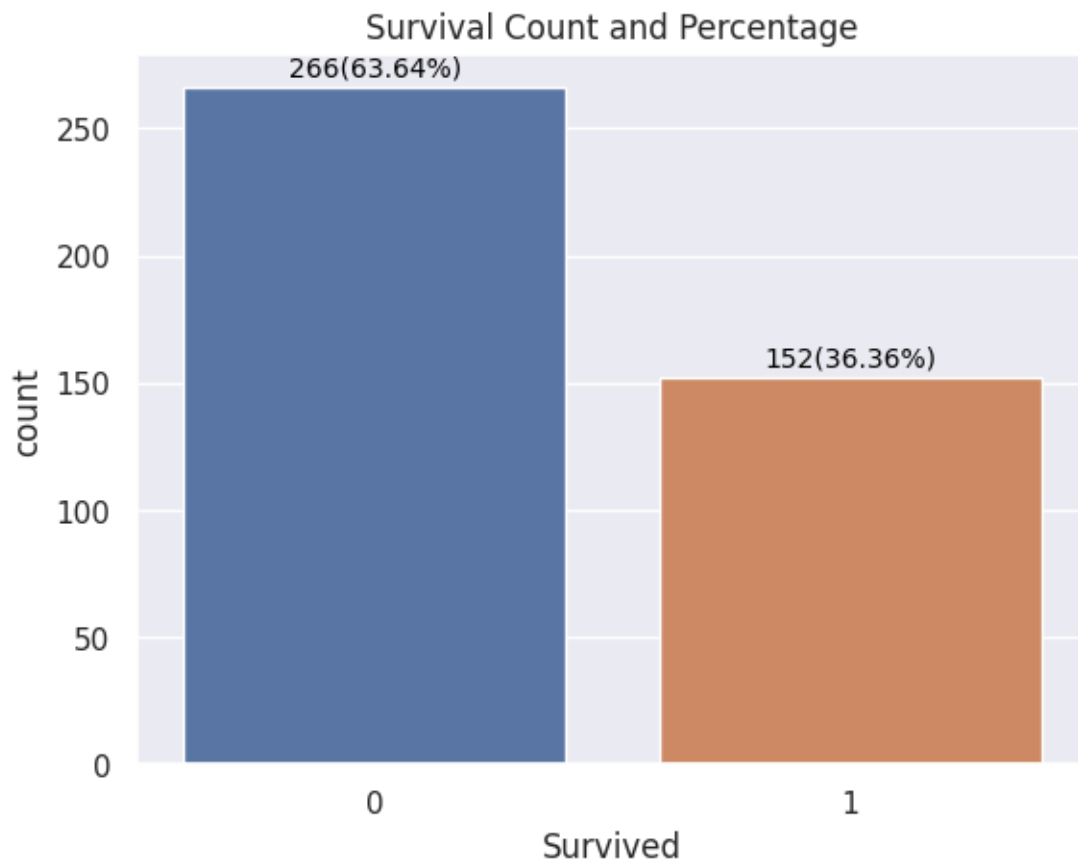


```
[239]: ax = sns.countplot(x='Survived', data=df)

total = len(df['Survived'])
for p in ax.patches:
    count = p.get_height()
    percentage = 100 * count / total
    x = p.get_x() + p.get_width() / 2
    y = p.get_height() + 5
    ax.annotate(f'{int(count)}({percentage:.2f}%)', (x, y), ha='center',
    ↪fontsize=10, color='black')

plt.title('Survival Count and Percentage')

plt.show()
```



```
[240]: class_counts = df['Pclass'].value_counts()

plt.pie(class_counts, labels=[f'Class {i} ({count})' for i, count in
    ↪class_counts.iteritems()], autopct='%1.1f%%', startangle=90,
    ↪wedgeprops=dict(width=0.4))

centre_circle = plt.Circle((0,0),0.70,fc='white')
fig = plt.gcf()
fig.gca().add_artist(centre_circle)


plt.axis('equal')

plt.title('Class Variation')

plt.show()
```

<ipython-input-240-f1b15da63cbd>:4: FutureWarning: iteritems is deprecated and will be removed in a future version. Use .items instead.

```
plt.pie(class_counts, labels=[f'Class {i} ({count})' for i, count in
class_counts.iteritems()], autopct='%1.1f%%', startangle=90,
wedgeprops=dict(width=0.4))
```



```
[241]: Age_Fare = ['Age', 'Fare']
for column in Age_Fare:
    plt.figure(figsize=(14,5))
    plt.subplot(1,2,1)
    ax = sns.boxplot(df[column])

    plt.subplot(1,2,2)
    ax = sns.distplot(df[column])
    plt.show()
```


<ipython-input-241-1fc3ed409439>:8: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
ax = sns.distplot(df[column])
```




<ipython-input-241-1fc3ed409439>:8: UserWarning:

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
ax = sns.distplot(df[column])
```




- In the 'Age' variable, there are some outliers present. The distribution of ages indicates that the majority of passengers fell within the age range of 20 to 40 years old.



- Regarding the 'Fare' variable, it exhibits several outliers, suggesting that there were passengers who paid significantly higher fares. Additionally, the distribution of fares is right-skewed, indicating that the majority of passengers paid lower fares, with a few paying much higher amounts.

```
[242]: sns.scatterplot(data=df, x='Age', y='Fare', hue='Survived')
plt.show()
```



The scatterplot shows that the passengers with high fares had more chance of survival

### To count & Visualise by Survived by people


Survived by Gender

```
[243]: ax = sns.countplot(x='Sex', data=df[df['Survived'] == 1])
total_survived = len(df[df['Survived'] == 1])

for p in ax.patches:
    count = p.get_height()
    x = p.get_x() + p.get_width() / 2
    y = p.get_height() + 0.05
    ax.annotate(f'{int(count)}', (x, y), ha='center', fontsize=10)
```

```
plt.title('Survivors by Gender')
plt.xlabel('Gender')
plt.ylabel('Count')

plt.show()
```



So, by the above graph you can see only females are alive in the titanic accident. By this we can confirm they followed “Womens and children first” if *an* accident occurs!!

Survived by Pclass

```
[244]: survived_by_pclass = df[df['Survived'] == 1]['Pclass'].value_counts()

plt.pie(survived_by_pclass, labels=[f'Pclass {i} ({count})' for i, count in
    ↪survived_by_pclass.iteritems()],
        autopct='%1.1f%%', startangle=90, pctdistance=0.85,
        ↪wedgeprops=dict(width=0.4))

centre_circle = plt.Circle((0, 0), 0.70, fc='white')
```

```

fig = plt.gcf()
fig.gca().add_artist(centre_circle)

plt.axis('equal')
plt.legend([f'Pclass {i}' for i in survived_by_pclass.index], title='Pclass',
           loc='center left', bbox_to_anchor=(1, 0, 0.5, 1))


for i, count in enumerate(survived_by_pclass):
    plt.text(0, -0.2 - i*0.15, f'{int(count)}', ha='center', fontsize=10)

plt.title('Survivors by Pclass')
plt.show()

```

<ipython-input-244-32ea0026edd4>:3: FutureWarning: iteritems is deprecated and will be removed in a future version. Use .items instead.

```
plt.pie(survived_by_pclass, labels=[f'Pclass {i} ({count})' for i, count in
survived_by_pclass.items()],
```



### Data value replacing with numbers

```
[245]: df.new=df.replace({'Sex':{'male':'0','female':'1'},'Embarked':{'S':'0','C':
           ↳'1','Q':'2'}})
```

<ipython-input-245-b067cfad61b6>:1: UserWarning: Pandas doesn't allow columns to be created via a new attribute name - see <https://pandas.pydata.org/pandas->

docs/stable/indexing.html#attribute-access

```
df.new=df.replace({'Sex':{'male':'0','female':'1'},'Embarked':{'S':'0','C':'1','Q':'2'}})
```

### Separating Features and Target

```
[246]: X=df.new.drop(columns=['PassengerId', 'Name', 'Ticket'])
      Y=df.new['Survived']
```

```
[247]: print(X)
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0	3	0	34.50000	0	0	7.8292	2
1	1	3	1	47.00000	1	0	7.0000	0
2	0	2	0	62.00000	0	0	9.6875	2
3	0	3	0	27.00000	0	0	8.6625	0
4	1	3	1	22.00000	1	1	12.2875	0
..	...	...	..	...	...	...	...	...
413	0	3	0	30.27259	0	0	8.0500	0
414	1	1	1	39.00000	0	0	108.9000	1
415	0	3	0	38.50000	0	0	7.2500	0
416	0	3	0	30.27259	0	0	8.0500	0
417	0	3	0	30.27259	1	1	22.3583	1

[418 rows x 8 columns]

```
[248]: print(Y)
```

```
0    0
1    1
2    0
3    0
4    1
..
413  0
414  1
415  0
416  0
417  0
Name: Survived, Length: 418, dtype: int64
```

### Training data & testing data

```
[249]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
      random_state=2)
```

```
[250]: print(X.shape, X_train.shape, X_test.shape)
```

(418, 8) (334, 8) (84, 8)

## Model Training

### Logistics Regression

```
[251]: model = LogisticRegression()
```

```
[252]: model.fit(X_train,Y_train)
```

```
[252]: LogisticRegression()
```

## Model Evaluation

### Accuracy Score

```
[253]: X_train_prediction=model.predict(X_train)
```

```
[254]: print(X_train_prediction)
```

```
[1 1 0 0 1 1 0 0 0 1 0 0 1 0 0 0 1 0 1 0 1 0 1 1 0 0 0 0 0 1 0 0 0 0 0 0 0
 1 1 1 0 0 0 1 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 1 0 1 0 1 1 1 0 1
 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 1 0 0 0 0 0 0 0 1 0 1 1 1 0 1 0 1 0
 1 1 0 0 0 0 1 1 0 1 0 0 1 1 0 1 0 0 0 0 0 0 1 0 0 1 0 0 1 0 0 1 0 1 1 0 0
 0 0 1 1 1 0 0 1 1 0 1 1 0 0 0 0 0 0 0 0 1 1 0 0 1 1 1 1 0 1 0 0 0 0 1 0 1 1
 1 0 1 0 0 0 1 0 0 0 1 0 1 0 0 0 0 0 0 0 1 1 1 1 0 0 0 0 1 0 0 1 0 0 1 0 0
 1 0 1 0 0 0 0 0 1 0 0 0 1 1 0 0 0 1 1 0 1 0 0 0 0 0 1 0 0 0 0 0 1 0 0 0 1
 0 1 1 1 1 0 0 0 1 1 0 0 1 0 1 1 0 0 0 0 1 0 0 0 0 0 1 0 0 1 1 0 1 1 0 0 0
 0 0 0 0 1 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 1 1 0 1 1 0 0 0 1 1 1
 1]
```

```
[255]: Training_data_accuracy= accuracy_score(Y_train,X_train_prediction)
print('Accuracy Score of training data: ',Training_data_accuracy)
```

Accuracy Score of training data: 1.0

```
[256]: X_test_prediction=model.predict(X_test)
Test_data_accuracy= accuracy_score(Y_test,X_test_prediction)
print('Accuracy Score of testing data: ',Test_data_accuracy)
```

Accuracy Score of testing data: 1.0

### Neighbour Classifier

```
[257]: from sklearn.neighbors import KNeighborsClassifier
```


```
knn = KNeighborsClassifier()
knn.fit(X_train,Y_train)
Y_pred_knn = knn.predict(X_test)
```

```
[258]: print('Accuracy: ', accuracy_score(Y_test, Y_pred_knn))
print('Precision: ', precision_score(Y_test, Y_pred_knn))
```

```
print('Recall: ', recall_score(Y_test, Y_pred_knn))
print('F1 Score: ', f1_score(Y_test, Y_pred_knn))
```

Accuracy: 0.7261904761904762  
Precision: 0.75  
Recall: 0.45454545454545453  
F1 Score: 0.5660377358490566

```
[260]: cm = confusion_matrix(Y_test, Y_pred_knn, labels=knn.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=knn.classes_)
disp.plot();
```



```
[261]: from sklearn.tree import DecisionTreeClassifier
```


```
tree = DecisionTreeClassifier()
tree.fit(X_train, Y_train)
y_pred_tree = tree.predict(X_test)
```

```
[262]: print('Accuracy: ', accuracy_score(Y_test, y_pred_tree))
print('Precision: ', precision_score(Y_test, y_pred_tree))
print('Recall: ', recall_score(Y_test, y_pred_tree))
```

```
print('F1 Score: ', f1_score(Y_test, y_pred_tree))
```

Accuracy: 1.0  
Precision: 1.0  
Recall: 1.0  
F1 Score: 1.0

```
[263]: cm = confusion_matrix(Y_test, y_pred_tree, labels=tree.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=tree.classes_)
disp.plot();
```



### Summary:

The treemap visualization indicates that the Decision Tree model achieved perfect accuracy in predicting passenger survival solely based on their gender. This is because all females survived the Titanic disaster, while all males did not.

### Key statistics:

- 36.36% of passengers survived the Titanic disaster.
- 63.64% of passengers did not survive.

### Gender-based survival:

- All females survived.
- None of the males survived.

**Class-based survival:**

- Passengers in Class 3 and 2 had a lower chance of survival.

**Model Performance:** Both the Logistic Regression and Decision Tree models achieved 100% accuracy.

**\*\*Thank You !!\*\***